



MapReduce

Kai Zhao





Review



Common Big Data Challenges

Volume:

- Too much data for processing → reduce processing time by using distributed & parallel computing
 - Performing word counting on 300 articles simultaneously
- Solution: **Distributed & Parallel Computing**

Parallel Computing

- Make use of **parallel architectures** (e.g. multi-core, multi-processor, clusters of machines, etc.) to **improve** computing performance
 - Using **multiple cores** to speedup video processing (Facetime HD!)
 - Using **multiple processors** in web servers
 - Using **GPUs** to increase rendering frame rate in games
 - Using **clusters** to run simulations
- Have a long history in High Performance Computing
 - Similar to streaming, getting **popular** in the Big Data era

Python's Core HOF

- `map()`: applies a function over an iterable to produce a new iterable
- `map(lambda x: int(x)+1, ['0', '1']) -> [1, 2]`
- `filter()`: return only values that satisfy a predicate in the new iterable
- `filter(lambda x: x<1, [0, 1]) -> [0]`
- `reduce()`: accumulate a sequence of values from left to right
- `reduce(lambda x, y: x+y, xrange(5), 100) -> 110`
- `sorted()`: return a new sorted list using a comparator function
- `sorted(xrange(5)) -> [0, 1, 2, 3, 4]`
- `sorted(xrange(5), cmp=lambda x, y: y-x) -> [4, 3, 2, 1, 0]`

Outline

- MapReduce
- Designing MapReduce Algorithms



MapReduce



Why is MapReduce?

- By definition, big data is too large to handle by conventional means. Sooner or later, you just **can't scale up** anymore
 - machines cannot grow too large...
- But we can **add more computers!**
- MapReduce paradigm allows us to **scale out** - even with commodity hardware
 - first proposed in the big data world by Google Jeffrey Dean and Sanjay Ghemawat.
 - MapReduce: Simplified Data Processing on Large Clusters. 2004.

What is MapReduce?

- A programming paradigm (for big data processing)
 - Data is **split** into distributable **chunks**
 - Perform **a series of transformations** on those chunks
 - Transformations are run **in parallel** on the chunks (data parallelism)
- MapReduce is scalable by **adding more machines** to process chunks
 - scaling out on commodity hardware
- The foundation for **Hadoop**, which is an implementation of MapReduce

What is MapReduce?

- A programming paradigm that process data in 2 phases/operations: `map()`, and then `reduce()`
 - instead of `map()`, `filter()`, `reduce()`, `sort()` like in the previous class, we are limited to only `map()` and `reduce()`
- In a nutshell — that's it:
 - Provide a `data collection` (separable records)
 - Apply a user-defined `map function` on `each` data record
 - Then `reduce` the mapped output with another user-defined `function`

What is MapReduce?

- MapReduce works on (key,value) pairs


INPUT: list of key-value pairs of (k_1, v_1)

MAP: $(k_1, v_1) \rightarrow [\text{list of } (k_2, v_2)]$

SHUFFLE: combine $(k_2, v_2) \rightarrow (k_2, [\text{list of } v_2])$

REDUCE: $(k_2, [\text{list of } v_2]) \rightarrow (k_3, v_3)$

OUTPUT: list of (k_3, v_3)



Fixed
pipeline

What is MapReduce?

- MapReduce works on (key,value) pairs

INPUT: list of key-value pairs of (k_1, v_1)



- MAP: $(k_1, v_1) \rightarrow [\text{list of } (k_2, v_2)]$

SHUFFLE: combine $(k_2, v_2) \rightarrow (k_2, [\text{list of } v_2])$

- REDUCE: $(k_2, [\text{list of } v_2]) \rightarrow (k_3, v_3)$

OUTPUT: list of (k_3, v_3)

User-defined functions
executed in parallel

Input

- Data must be separable into records
 - Lines of text
 - Rows of tables
 - CSV: yes — JSON, XML: no
- Key/value pairs
 - Key = line number, record index
 - Value = text string, row data
- Keys are mostly ignored in many cases (e.g. just passing text lines)

Map Phase

- Transform each input record using a user-defined **function**
 - **$(k1, v1) \rightarrow [(k2, v2), \dots]$** — could be an empty list
- **one-to-many** relationship
 - vs. one-to-one of Python's `map()` higher order function
 - Python's `filter()` is part of the Map phase
- Mainly process by streaming data chunks
 - **$[(k1, v1)] \rightarrow [(k2, v2), \dots]$**
 - generators \rightarrow generators

Shuffle and Sort Phase

- For each key **k2**, collect all **v2** from outputs of all map processes — shuffling into **(k2, [v2])**
 - Done by distributed partitioning and local sort
- Outputs are **(k2, [v2])** being distributed to “reducers”
- **[(k2,v2)] → (k2, [v2])**
- **(k2,v2)** are intermediate key/value pairs

Reduce Phase



- Transform the output of the Shuffle phase using a user-defined function
 - $(k2, [v2]) \rightarrow (k3, v3)$
- Equivalent to `reduceByKey()` ~ `groupByKey()` x `reduce()`
 - vs. Python's `reduce()`
- • Also process in streaming data chunks
 - $[(k2, [v2])] \rightarrow [(k3, v3)]$
- This phase is optional, can be omitted.

The Classic Example: WordCount

- The “Hello World” of MapReduce
- Given a set of documents, compute the frequency of each word
- INPUT (k1, v1): (document: line number, text line)
- OUTPUT (k3, v3): (word, frequency)
- INTERMEDIATE (k2, v2): (word, 1)
 - Transform each line to a list of words and their frequencies (=1)
 - Combine all tuples by words, and add all the frequency together


The Classic Example: WordCount

- The “Hello World” of MapReduce
- Given a set of documents, compute the frequency of each word



MapReduce Dataflow

Intermediate (key,value) pairs



The Shuffle

When to use MapReduce?

- Data parallelism problems (“embarrassingly parallel”)
 - Word count
 - Distributed grep
 - Reverse index
 - Document OCR (Optical character recognition)
- Data must be splittable into chunks and records



Designing MapReduce Algorithms



Designing MapReduce Algorithms

- Key decision: What should be done by **map**, and what by **reduce**?
 - **map** can do something to each **individual** key-value pair, but it can't look at other key-value pairs
- Example: Filtering out key-value pairs we don't need
 - **map** can **emit** more than one intermediate key-value pair for each incoming key-value pair
- Example: Incoming data is text, map produces (word,1) for each word

Designing MapReduce Algorithms

- Key decision: What should be done by map, and what by reduce?
 - **reduce** can **aggregate** data; it can look at multiple values, as long as map has mapped them to the **same** (intermediate) **key**
 - Example: Count the number of words, add up the total cost, ...
- Need to get the intermediate format right!
 - If reduce needs to look at several values together, map must emit them using the same key!

Filtering with MapReduce


- Goal: find lines/files/tuples with a particular characteristic
- Examples:
 - grep Web logs for requests to a specific domain
 - find in the Web logs the hostnames accessed by a particular IP
 - locate all the files that contain the words 'Apple' and 'Jobs'
- Generally: map does most of the work, utilizing one-to-many relationships to discard unqualified records. reduce may simply be the identity

Aggregation with MapReduce

- Goal: compute the **maximum, the sum, the average**, ..., over a set of values
- Examples:
 - Count the number of requests to a website
 - Find the most popular domain
 - Average the number of requests per page per Web site
- Often: map may be simple or the identity, reduce does the aggregation

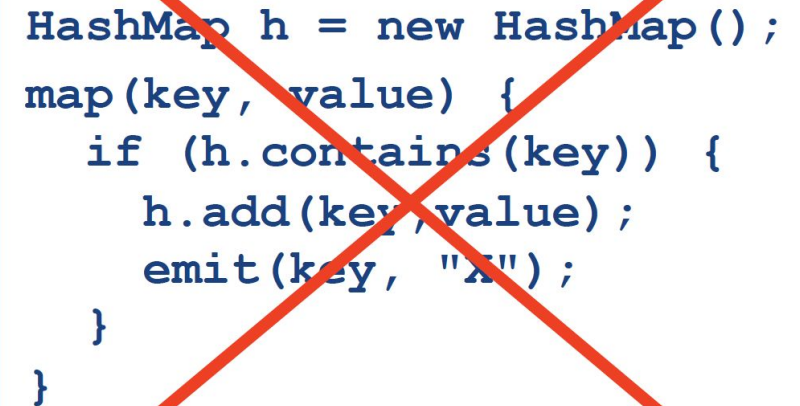
Combiners

- Certain functions can be decomposed into **partial steps**:
 - Can take counts of two **sub-partitions**, sum them up to get a complete count for the partition
 - Can take maxes of two sub-partitions, max them to get a complete max for the partition
- Multiple map jobs on the same machine may write to the same reduce key



Common Mistakes to Avoid

- Mapper and reducer should be **stateless**
-
- Don't use static variables - after map +
- reduce return, they should **remember**
- **nothing** about the processed data!
-
- Reason: No guarantees about which
- key-value pairs will be processed by
- which workers!



```
HashMap h = new HashMap();  
map(key, value) {  
    if (h.contains(key)) {  
        h.add(key, value);  
        emit(key, "X");  
    }  
}
```

Wrong!

Common Mistakes to Avoid

- Mapper must not map too much data to the same key
- In particular, don't map everything to **the same key!!** Otherwise the reduce worker will be **overwhelmed!**
- It's okay if some reduce workers have more work than others

```
map(key, value) {  
    emit("FOO", key + " " + value);  
}
```

Wrong!